
MODULE *UpdateCluster*

EXTENDS *Integers, FiniteSets*

CONSTANTS

_Requests, the requests sent by the user
_Workers, the pool of workers
NULL

VARIABLES

lastVOK, last successfully applied version
toApply, the version to apply (last request that passed the initial tests)
cluster, cluster state
requests, the state of all requests
workers, the state of all workers
clusterUpdating damn, I use a lock...

VARIABLES

these variables are *tla* + details
confOK, are we able to get a valid conf?
reqCounter just to keep track of the order of submissions

vars $\triangleq \langle \textit{confOK}, \textit{reqCounter}, \textit{lastVOK}, \textit{toApply}, \textit{cluster}, \textit{requests}, \textit{workers}, \textit{clusterUpdating} \rangle$

TypeInvariants \triangleq

$\wedge \textit{confOK} \in \text{BOOLEAN}$ won't change for a specific behavior
 $\wedge \textit{clusterUpdating} \in \text{BOOLEAN}$
 $\wedge \textit{cluster.st} \in \{$
 "idle",
 "starting",
 "partial",
 "failed"
 }
 $\wedge \forall r \in \textit{_Requests} : \textit{requests}[r].\textit{st} \in \{$
 "waiting", the request (req) hasn't been submitted yet
 "submitted", req has been submitted
 "rejected", req has been rejected (auth problem)
 "valid" auth etc passed
 }
 $\wedge \forall w \in \textit{_Workers} : \textit{workers}[w].\textit{st} \in \{$
 "waiting",
 "starting",
 "working"
 }

Initial State

$Init \triangleq$
 $\wedge requests = [r \in _Requests \mapsto [st \mapsto \text{"waiting"}, v \mapsto NULL]]$
 $\wedge workers = [w \in _Workers \mapsto [st \mapsto \text{"waiting"}, v \mapsto NULL]]$
 $\wedge cluster = [v \mapsto 0, st \mapsto \text{"idle"}]$
 $\wedge lastVOK = 0$
 $\wedge reqCounter = 0$
 $\wedge toApply = 0$
 $\wedge confOK \in \text{BOOLEAN}$
 $\wedge clusterUpdating = \text{FALSE}$

Actions

request actions
 $Submit(r) \triangleq$
 $\text{update request received}$
 $LET \ newV \triangleq reqCounter + 1 \text{ IN}$
 $\wedge requests[r].st = \text{"waiting"}$
 $\wedge reqCounter' = newV$
 $\wedge requests' = [requests \text{ EXCEPT } ![r].st = \text{"submitted"}, ![r].v = newV]$
 $\wedge \text{UNCHANGED } \langle confOK, lastVOK, toApply, cluster, workers, clusterUpdating \rangle$

$PushToPending(r) \triangleq$
 $\text{the request is pushed}$
 $\wedge requests[r].st = \text{"submitted"}$
 $\wedge \text{IF } toApply < requests[r].v$
 $\quad \text{THEN } \wedge requests' = [requests \text{ EXCEPT } ![r].st = \text{"valid"}]$
 $\quad \wedge toApply' = requests[r].v$
 $\quad \wedge \text{UNCHANGED } \langle confOK, reqCounter, lastVOK, cluster, workers, clusterUpdating \rangle$
 $\quad \text{ELSE } \wedge requests' = [requests \text{ EXCEPT } ![r].st = \text{"rejected"}]$
 $\quad \wedge \text{UNCHANGED } \langle confOK, reqCounter, lastVOK, toApply, cluster, workers, clusterUpdating \rangle$

worker action
 $SpawnWorker(w) \triangleq$
 $\text{spawns a new worker}$
 $\wedge toApply \neq lastVOK$
 $\wedge workers[w].st = \text{"waiting"}$
 $\wedge clusterUpdating = \text{FALSE}$
 $\wedge \vee cluster.st = \text{"idle"}$
 $\quad \vee cluster.st = \text{"failed"}$
 $\wedge workers' = [workers \text{ EXCEPT } ![w].v = toApply, ![w].st = \text{"starting"}]$
 $\wedge clusterUpdating' = \text{TRUE}$
 $\wedge \text{UNCHANGED } \langle confOK, reqCounter, lastVOK, toApply, requests, cluster \rangle$

$ApplyStart(w) \triangleq$
 the cluster starts to be modified
 $\wedge workers[w].st = \text{"starting"}$
 $\wedge \text{IF } workers[w].v \geq toApply$
 THEN
 IF $confOK$
 THEN
 $\wedge cluster' = [v \mapsto workers[w].v, st \mapsto \text{"partial"}]$
 $\wedge workers' = [workers \text{ EXCEPT } ![w].st = \text{"working"}]$
 $\wedge \text{UNCHANGED } \langle confOK, reqCounter, lastVOK, toApply, requests, clusterUpdating \rangle$
 ELSE
 $\wedge clusterUpdating' = \text{FALSE}$
 $\wedge workers' = [workers \text{ EXCEPT } ![w].st = \text{"waiting"}, ![w].v = NULL]$
 $\wedge \text{UNCHANGED } \langle confOK, reqCounter, lastVOK, toApply, cluster, requests \rangle$
 ELSE a new version has been submitted, no need to apply this one
 $\wedge workers' = [workers \text{ EXCEPT } ![w].st = \text{"waiting"}, ![w].v = NULL]$
 $\wedge clusterUpdating' = \text{FALSE}$
 $\wedge \text{UNCHANGED } \langle confOK, reqCounter, lastVOK, toApply, cluster, requests \rangle$

$RollbackVersion \triangleq$
 to differentiate it from the original last VOK (in realworld, could be $conf + \text{timestamp}$)
 $lastVOK + 10$

$ApplyFinish(w) \triangleq$
 the cluster update finishes
 $\wedge workers[w].st = \text{"working"}$
 $\wedge clusterUpdating' = \text{FALSE}$
 $\wedge \exists ok \in \text{BOOLEAN} :$
 IF $ok \vee workers[w].v = RollbackVersion$ rollback always works
 THEN
 $\wedge cluster' = [cluster \text{ EXCEPT } !.st = \text{"idle"}]$
 $\wedge lastVOK' = workers[w].v$
 $\wedge workers' = [workers \text{ EXCEPT } ![w].st = \text{"waiting"}, ![w].v = NULL]$
 $\wedge \text{UNCHANGED } \langle confOK, reqCounter, toApply, requests \rangle$
 ELSE
 $\wedge cluster' = [cluster \text{ EXCEPT } !.st = \text{"failed"}]$
 $\wedge workers' = [workers \text{ EXCEPT } ![w].st = \text{"waiting"}, ![w].v = NULL]$
 $\wedge \text{IF } workers[w].v < toApply$
 THEN a newer version has been submitted
 $\wedge \text{UNCHANGED } \langle confOK, reqCounter, lastVOK, toApply, requests \rangle$
 ELSE let's trigger a rollback
 $\wedge toApply' = RollbackVersion$
 $\wedge \text{UNCHANGED } \langle confOK, reqCounter, lastVOK, requests \rangle$

Spec

$Next \triangleq$

$$\begin{aligned} & \vee \exists r \in _Requests : \\ & \quad \vee Submit(r) \\ & \quad \vee PushToPending(r) \\ & \vee \exists w \in _Workers : \\ & \quad \vee SpawnWorker(w) \\ & \quad \vee ApplyStart(w) \\ & \quad \vee ApplyFinish(w) \end{aligned}$$

$Fairness \triangleq \forall r \in _Requests, w \in _Workers :$

$$\begin{aligned} & \wedge WF_{vars}(Submit(r)) \\ & \wedge WF_{vars}(PushToPending(r)) \\ & \wedge WF_{vars}(SpawnWorker(w)) \\ & \wedge WF_{vars}(ApplyStart(w)) \\ & \wedge WF_{vars}(ApplyFinish(w)) \end{aligned}$$

$Spec \triangleq$

$$\begin{aligned} & \wedge Init \\ & \wedge \Box[Next]_{vars} \\ & \wedge Fairness \end{aligned}$$

Expectations

$NoConcurrentUpdate \triangleq$

$$\Box(Cardinality(\{w \in DOMAIN\ workers : workers[w].st = \text{"working"}\}) < 2)$$

$NoPartialUpdateTermination \triangleq$

we don't want the cluster to end up in a partially update state
 $\Diamond\Box(cluster.st = \text{"idle"})$

$EveryReqIsProcessed \triangleq$

$$\Diamond\Box(\neg\exists r \in _Requests : requests[r].st = \text{"waiting"})$$

THEOREM $Spec \Rightarrow \Box(TypeInvariants)$

THEOREM $Spec \Rightarrow NoConcurrentUpdate$

THEOREM $Spec \Rightarrow NoPartialUpdateTermination$

THEOREM $Spec \Rightarrow EveryReqIsProcessed$